



COP4635: Systems & Networks II

Network Programming:
Sockets & Client Server



Overview

- Sockets
 - communication endpoints for processes
 - can be on different machines
 - can be used over network
- Client - Server
 - client is a process sending requests to a server
 - server is a continuously running process responding to client requests
 - client – server use connection-oriented service

LD2 Network programming

COP4635

2



Socket-Based Communication

- Socket
 - Communication endpoint
 - Commonly used for network communication
 - Composed of three parts
 - **IP Address**: Identifies machine
 - **Port**: Identifies process on machine
 - **Protocol**: Specifies communication parameters, order, etc.


LD2 Network programming

COP4635

3



Network Communication



LD2 Network programming

COP4635

4



Relevant Data Types

The following data types are needed to build sockets:

u_char	unsigned char (8-bit)
u_short	unsigned short (16-bit)
u_long	unsigned long (32-bit)

LD2 Network programming

COP4635

5



Internet Address

- Specifies machine on Internet
 - Assume we are using IPv4 (not IPv6)
 - 32-bit integer number
 - Usually written as dotted quad
 - Type is u_long (unsigned long)
 - Wrapped in structure

```
struct in_addr
{
    u_long    s_addr;
};
```

LD2 Network programming

COP4635

6



Port Numbers

Logical not physical
16-bit integer
(short)

0-1023	Well-known (reserved)
1024-49151	Registered (discouraged)
49152-65535	Dynamic/Private

7	echo	25	smtp	113	ident
11	systat	37	time	119	nnntp
21	ftp	79	finger	135	RPC
22	ssh	80	http	143	imap
23	telnet	110	pop3	443	https

LD2 Network programming

COP4635

7



Protocol

- Specifies
 - Format of messages
 - Headers
 - Other communication parameters
 - A few examples:
 - PF_INET
 - PF_INET6
 - PF_BLUETOOTH
 - PF_APPLETALK

LD2 Network programming

COP4635

8



SOCKADDR

```
struct sockaddr_in
{
    u_char          sin_len;
                        // normally not used
    u_short         sin_family;
    u_short         sin_port;
    struct in_addr  sin_addr;
    char            sin_zero[8];
                        // normally not used
};
```

LD2 Network programming

COP4635

9



Socket Types

- TCP
- Connection-oriented, reliable
 - Also called *stream*
- UDP
- Connectionless, unreliable
 - Also called *datagram*
- RAW
- Bypasses niceties of above
 - Functions at IP level


LD2 Network programming

COP4635

10



Protocol Layers



LD2 Network programming

COP4635

11



Potential Problem


- All machines are not created equal.
2- & 4-byte values can be stored differently.
- ➡ **Endianness** (byte order) can differ.

The very term *big-endian* comes from Jonathan Swift's satiric novel *Gulliver's Travels*, where tensions are described in Lilliput and Blefuscu: whereas royal edict in Lilliput requires cracking open one's soft-boiled egg at the small end, inhabitants of the rival kingdom of Blefuscu crack theirs at the big end (hence the name *Big-endians*).

LD2 Network programming

COP4635

12



Byte Ordering

143
8F

88
58

65
41

220
DC

MS byte

LS byte

L-E

A

A+1

A+2

A+3

B-E

A

A+1

A+2

A+3


x86
VAX
PDP-11

Motorola 68000
PowerPC
SPARC
networks

LO2 Network programming

COP4635

13



The Problem

MS byte

LS byte

143

88

65

220

10001111

01011000

01000001

11011100

L-E

A

A+1

A+2

A+3

B-E

A

A+1

A+2

A+3

2404925916

3695270031

MS byte

LS byte

220

65

88

143

11011100

01000001


01011000

10001111

LO2 Network programming

COP4635

14



Communication

Machine A

- Big Endian
- Send 2-byte decimal value 821
- In binary, that's (00000011 00110101)


Machine B

- Little Endian
- Receive binary (00110101 00000011)
- In decimal, that's 13571

LO2 Network programming

COP4635

15




Network Byte Order

- Network bytes order is big-endian
- Special functions to convert to/from NBO
 - htons host to network short
 - ntohs network to host short
 - htonl host to network long
 - ntohl network to host long
- Trick
 - B-E machine: functions do nothing
 - L-E machine: functions reverse bytes
 - Always use on data
 - Don't worry about endianness of machine

LO2 Network programming

COP4635

16



Address Transformation

Lots of functions to manipulate IP Addresses.

Refer to man pages for details and include files.


```
int inet_aton( char* cp,
               struct in_addr *inp)
    Converts dotted-quad (cp) to 32-bit address in network byte
    order (inp)

char *inet_ntoa( struct in_addr in)
    Converts 32-bit address in network byte order (in) to dotted-
    quad (return pointer)
```

LO2 Network programming

COP4635

17



Host Information

Given hostname, we need to lookup the IP address:

```
struct hostent * gethostbyname( char* name)

struct hostent
{
    char    *h_name;
    char    **h_aliases;
    int     h_addrtype; /* PF_INET */
    int     h_length;   /* 4 */
    char    **h_addr_list;
```

LO2 Network programming

COP4635

18



Creating a Socket

Access similar to pipe or file
Through file descriptor table
Called *socket descriptor* in this case

```
int socket(
    int domain,    /* AF_INET */
    int type,      /* SOCK_DGRAM, SOCK_STREAM */
    int protocol /* 0 */
);
```

Return value of -1 indicates an error

LD2 Network programming

COP4635

19



Putting It All Together

1. Create a socket
2. Get information about *dest* host.
3. Fill in *dest addr* structure.
4. Connect to remote host.
5. Send /receive data to *dest*.
6. Close connection.

LD2 Network programming

COP4635

20



Create a socket

```
socketFD = socket( AF_INET, SOCK_DGRAM, 0 );
if ( socketFD < 0 ) {
    perror( "sendUDP:socket" );
    return -1;
}
```

LD2 Network programming

COP4635

21



Get Dest Host Info

Obtaining information about a remote host:

```
hostptr = gethostbyname( argv[1] );
if ( hostptr == NULL ) {
    perror( "sendUDP:gethostbyname" );
    return -1;
}
```

Function makes a call to DNS and returns a struct called *hostent*.

LD2 Network programming

COP4635

22



Fill In Dest Addr

```
memset( &destaddr, 0,
        sizeof(struct sockaddr_in) );
destaddr.sin_family = AF_INET;
memcpy( &destaddr.sin_addr,
        hostptr->h_addr,
        hostptr->h_length );
destaddr.sin_port =
    htons( (u_short)atoi(argv[2]) );
```

LD2 Network programming

COP4635

23



Send Message

```
fgets( buffer, 256, stdin );
sendlen = strlen( buffer ) - 1;
bytes = sendto(
    socketFD,
    buffer,
    (size_t)sendlen,
    0,
    (struct sockaddr *)&destaddr,
    (socklen_t)sizeof(destaddr)
);
```

LD2 Network programming

COP4635

24



Receiver Extra Step

- Receiver must have port number.
- Delivery:
 - Packet → Machine (host) → OS → Process
- The `bind()` system call provides “name” for socket.
- Can send without a “named socket”, but must have “named socket” to receive.
 - Similar to U.S. Postal Service

L02 Network programming

COP4635

25



Using bind

```
int bind(
    int sockfd,
    struct sockaddr *my_addr,
    socklen_t addrlen
);
```

L02 Network programming

COP4635

26



What Is sockaddr?

```
struct sockaddr_in
{
    u_char      sin_len;
    u_short     sin_family;
    u_short     sin_port;
    struct in_addr sin_addr;
    char        sin_zero[8];
};
```

L02 Network programming

COP4635

27



So...

`sendto()` → `recvfrom()`

Sender needs to send to a socket

- IP Address
- Protocol Family
- Port number

Receiver needs to establish socket

- Needs to establish port number
- How?

L02 Network programming

COP4635

28



Receiver

If 1st action is receive,

Must establish port number

And announce port number

Why?

So sender knows where to send

Otherwise

Receiver is “out there” somewhere

L02 Network programming

COP4635

29



Two Ways - Static

Static binding

- User provides port number; some options
- Hardcoded in program

```
DA.sin_port = htons( 51664 );
```

- Provided on command-line

```
DA.sin_port =
    htons((u_short)atoi(argv[1]));
```

- Provided via user input

```
DA.sin_port =
    htons((u_short)atoi(inString));
```

L02 Network programming

COP4635

30



Two Ways - Dynamic

Dynamic binding

- **Kernel** provides port number
- Send special “you pick” to kernel
- Value of 0 (zero) is indicator

```
DA.sin_port = 0;
```

- Kernel picks next available port number
- Stores port number internally

L02 Network programming

COP4635

31



Example

```
int bind(
    int sockfd,
    struct sockaddr *my_addr,
    socklen_t addrlen
);
```

The `bind()`
selects next
available port

```
int getsockname(
    int sockfd,
    struct sockaddr *my_addr,
    socklen_t *addrlen
);
```

The
`getsockname()`
returns the *name* of
the socket

L02 Network programming

COP4635

32



Announcing Socket Name

```
printf("name: %s\n",
    hostptr->h_name);

printf("addr: %s\n",
    inet_ntoa(srcAddr.sin_addr));

printf("port: %d\n",
    ntohs(srcAddr.sin_port));
```

L02 Network programming

COP4635

33



Executions

```
> ./recvUDP
name: sse-250-061.cs.uwf.edu
addr: 143.88.64.125
port: 54464
> ./recvUDP
name: sse-250-061.cs.uwf.edu
addr: 143.88.64.125
port: 54208
>
```

L02 Network programming

COP4635

34



What If...

Do you need a port number to send?

Yes

Where do you get it?

Dynamically – from kernel

How?

`sendto()` allocates port if needed

This is second form of dynamic binding

Try using `sendto()` then `getsockname()`

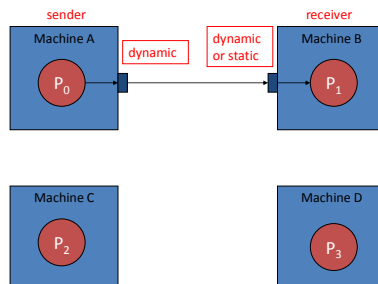
L02 Network programming

COP4635

35



Big Picture - A



L02 Network programming

COP4635

36



Logistics for Project Work

Sit at ONE machine (machine A)

Open a terminal

- Terminal displays on A
- Commands run on A

Open a second terminal

- Terminal displays on A
- Commands run on A
- Use `ssh` to connect to machine B
- Terminal displays on A
- Commands run on B

LD2 Network programming

COP4635

37



What About My Files?

Multiplatform Lab uses NFS

- Network File System
- Files stored on “file server”
- Files available on all machines
- Save on A, new file also visible on B

Compile on A, execute on B

LD2 Network programming

COP4635

38



On sse-250-061

display

sse-250-061

execute

sse-250-061

```
> ./recvUDP
name: sse-250-061.cs.uwf.edu
addr: 143.88.64.206
port: 13925
got [hello]
```

sse-250-061

sse-250-05e

```
> ssh sse-250-05e.cs.uwf.edu
> ./sendUDP sse-250-061.cs.uwf.edu 13925
Send this> hello
sent 5 bytes
>
```

LD2 Network programming

COP4635

39



Network Communication

Requires endpoints (sockets)

- Protocol (PF_INET)
- IP Address (32-bit or dotted quad)
- Port Number (16-bit; static or dynamic)

UDP

- Sends *datagrams*
- May arrive; may be lost
- May arrive in different order than sent

LD2 Network programming

COP4635

40



TCP

TCP is a protocol designed for the reliable transfer of bytes from one process to another over a network

How?

- Sequence numbers
- ACK
- Timers (what if small file?)

LD2 Network programming

COP4635

41



TCP Overview

Designed for client/server model

Server is at well-known place

Usually only one server

Clients (multiple) are anywhere

Client contacts server for *service*

Client and server communicate to facilitate service

Client and server disconnect

Server continues to wait for next client contact

LD2 Network programming

COP4635

42



TCP Notes

- Client and Server have “1-to-1” connection
- No other process can use connection
- Connection is bidirectional
 - Client can send to server
 - Server can send to client
- Requires buffers, segments, handshakes, ...



UDP vs TCP

- All (sender, receiver, client, server)
- `socket()`
 - `gethostbyname(who)`
- TCP must create client/server connection
- | <u>Client</u> | <u>Server</u> |
|------------------------|-----------------------|
| | <code>bind()</code> |
| | <code>listen()</code> |
| <code>connect()</code> | <code>accept()</code> |



TCP Server

- `bind()`
- Makes receipt before send possible
- `listen()`
- Expresses willingness to accept incoming connections
 - Sets connection parameters (queue limit, etc.) for incoming connections
- `accept()`
- Accepts a connection on a listening socket
 - Returns file descriptor for newly created 1-to-1 socket




TCP Handshake


- Done by `connect()` / `accept()`
- Client
 - Sends “CONN REQ” to server socket
- Server
 - Accepts request
 - Creates new socket (dynamic port binding)
 - Randomly generates 1st S→C sequence number
 - Sends new socket & 1st S→C sequence number to client (“CONN ACK”)
- Client
 - Accepts message
 - Randomly generates 1st C→S sequence number
 - Sends 1st C→S sequence number to server (“CONNECT”)



Client to Server




One-to-One Communication






Next Client



Still One-to-One



TCP Client Calls

```
socket()
gethostbyname( server )
connect()
send() / recv()
close()
```



TCP Server Calls

```
socket()
gethostbyname( self )
listen()
accept()
send() / recv()
close()
```




Normal Server Operation

```
listen()
while ( 1 ) {
    newFD = accept()
    pthread_create()
}
```

- Child thread
- Handles request
 - Communicates with client
 - Terminates when service complete
 - May block for I/O



Threaded Server





UDP

Sender	Receiver
<pre>socket() gethostbyname(dest) sendto() close()</pre>	<pre>socket() gethostbyname(dest) bind() recvfrom() close()</pre>



L02 Network programming

COP4635

55



FTP via UDP

Datagram is limited in size (i.e., 10K)

How can we transfer large file (i.e., 100K)?

UDP Sender

- Partition file into 10 datagrams
- Send datagram 1, then datagram 2, ...

UDP Receiver

- Receive all datagrams
- Reconstruct file from datagrams

L02 Network programming

COP4635

56



FTP/UDP Problem

Some datagrams could be lost

1 2 5 6 8 9 10

Datagrams could arrive out of order

1 2 5 8 10 3 9 7 4 6

Or Both

1 5 8 10 7 4 6

What now?

- We didn't receive file as it was sent

L02 Network programming

COP4635

57



UDP Solution?

Add a number to each datagram

- Sequence number (start at 0)
- Receiver can reorder datagrams
- Receiver can ask for lost datagram to be resent
- Receiver can send acknowledgement (ACK) of receipt for datagrams

L02 Network programming

COP4635

58



Creating Streaming Sockets in Java

- A socket is created using the `java.net.Socket`:

```
Socket socket = new Socket(host, port);
```

 host is a string specifying a host; can be localhost
 port is the port number on which the host is contacted
- Example:

```
Socket socket = new
  java.net.Socket(localhost, 8080);
```

L02 Network programming

COP4635

59



Creating Datagram Sockets in Java

- A datagram socket is created using

```
java.net.DatagramSocket:
```

```
DatagramSocket dsocket =
  new DatagramSocket(port);
```

 port is the UDP port on which packets will be received;
 the datagram socket is bound to host's Internet address
- Example:

```
DatagramSocket dsocket =
  java.net.DatagramSocket(8080);
```

L02 Network programming

COP4635

60



Streaming Sockets: Reading/Writing Data

- Using a socket, a data stream for reading and writing messages can be created.
- IO from a socket stream is like file IO:

```
// the buffer that stores read data
byte[] buffer = new byte[1024];

// creating an input stream; an error occurs if the
// socket is not open or not connected
InputStream iStream = socket.getInputStream();
OutputStream oStream = socket.getOutputStream();

// sending data stored in the buffer
oStream.write(buffer)

// reading data stored in the buffer
int numBytesRead = iStream.read(buffer)
```

L02 Network programming

COP4635

61



Streaming Sockets: Terminating a Connection

- To end a connection, either side of the end can call `close()`.

For example:

```
socket.close();
```

- If one end closes the socket and the other end does a
 - `read()`, the function will throw an IO exception,
 - `write()`, the function will throw an IO exception.

L02 Network programming

COP4635

62



Server Socket

- Server must listen for incoming connections.
 - server socket waits for a client connection
 - when a client connects to a server, the server creates a new socket, to communicate with the connecting host.
 - new socket communicates with connecting client on a port other than the server's listening port.
- Server sockets in Java (`java.net.ServerSocket`):
 - A server socket is created as follows:


```
ServerSocket ssocket = new ServerSocket(port);
```

 port is the port on which the server listens for connections
 - Example:


```
ServerSocket ssocket = ServerSocket(8080);
```

L02 Network programming

COP4635

63



Designing a Network Protocol

- Server and client respond to messages based on a well-defined protocol.
- Protocol defines actions and data that a client requests from a server.
- Client and server compile and parse message based on protocol.
 - protocol may specify type and order of data in message
 - meta data may describe data
- Message can be in different formats including
 - textual: all data are in text format
 - binary: all data are in binary format
 - mixed: data can be binary or textual

L02 Network programming

COP4635

64



Network Protocol Examples

- Online banking that supports users perform various banking transactions:


```
<action>operation</action><value>aValue</value>
```
- Browser sends message to a server to retrieve a Web page.

(see lab presentation)

L02 Network programming

COP4635

65



Server-Client: Online Banking Example

Client sending a message to the server:

```
private String sendMessage(String message) throws Exception {
    Socket socket;
    InputStream iStream;
    OutputStream oStream;
    try {
        // creates socket to communicate with server
        socket = new Socket(_host, _port);

        // access data streams from socket
        iStream = socket.getInputStream();
        oStream = socket.getOutputStream();
    }
    catch (Exception xcp) {
        throw new Exception ("connection error: unable to connect to the server");
    }
}
```

L02 Network programming

COP4635

66



Server-Client: Online Banking Example

Client sending a message to the server:

```
String response;
try {
    PrintWriter printWriter = new PrintWriter(oStream);
    Scanner scanner = new Scanner(iStream);

    // sending message
    printWriter.println(message);
    printWriter.flush();

    // receiving response
    response = scanner.nextLine();
} catch (Exception xcp) {
    socket.close();
    throw new Exception ("unable to communicate with server");
}
// close connection to server
socket.close();

return response;
}
```

PrintWriter and Scanner aid in writing and reading text from a socket.

LO2 Network programming

COP4635

67



Server-Client: Online Banking Example

Sample protocol for banking transaction:

Client sends a message to the server to create an account:

create|name|deposit
name: name of the account owner
deposit: initial deposit of the account

Example:

create|joe smith|300

Server sends message confirming account creation:

account_number
account_number: the number of the account

Example:

5

LO2 Network programming

COP4635

68



Server-Client: Online Banking Example

Server socket waiting for incoming connection:

```
ServerSocket serverSocket;
try {
    serverSocket = new ServerSocket(_port);
} catch (Exception xcp) {
    throw new Exception ("unable to start server on port " + _port);
}
while (true) {
    Socket socket = serverSocket.accept();
    try {
        // get communication channels of socket
        InputStream iStream = socket.getInputStream();
        OutputStream oStream = socket.getOutputStream();

        // start a request handler in a separate thread
        RequestHandler requestHandler = new RequestHandler(iStream, oStream);
        new Thread(requestHandler).start();

        // do nothing; client has dropped connection
    } catch (Exception xcp) {
        // do nothing; client has dropped connection
    }
}
```

Socket is accepting a connection to a host.

Handling request from client in a separate thread.

LO2 Network programming

COP4635

69



Server-Client: Online Banking Example

Handling a client's request:

```
public void run() {
    // read data from
    String requestMessage = _scanner.nextLine();
    String response = handleRequest(requestMessage);

    // send response back to host
    sendMessage(response);

    // close communication channels
    _printWriter.close();
    _scanner.close();
}
```

Read request from client and handle it.

Sends computed response.

Closes communication.

LO2 Network programming

COP4635

70



Server-Client: Online Banking Example

Creating a bank account using information submitted by client:

```
public int createAccount(String name, float initialBalance) throws Exception {
    int accountNumber = _bankData.getNextAccountNumber();
    BankAccount bankAccount = new BankAccount(accountNumber, name, initialBalance);
    _bankData.write(bankAccount);

    return accountNumber;
}
```

Sends message back to client:

```
public void sendMessage(String message) {
    try {
        _printWriter.println(message);
        _printWriter.flush();
    } catch (Exception xcp) {
        System.err.println("error occurred sending message to host: " + message);
    }
}
```

LO2 Network programming

COP4635

71



Server-Client: Online Banking Example

- Server accepts incoming connections continuously.
 - `accept()` blocks server until a new connection is requested
 - requests from clients are handled in separate threads
- Server terminates when program is shutdown.
 - more graceful shutdown could be done through a special server administration console
 - a client sends a shutdown message to the server
 - shutdown message causes server to close any open resources and to terminate

LO2 Network programming

COP4635

72